

Inference Latency on Raspberry Pi 4 ARM Cortex-A72
(lower is better — 100ms = real-time threshold)

